



Politecnico
di Torino

AI  ML

Exercises

Advanced Machine Learning

Teaching Assistant:

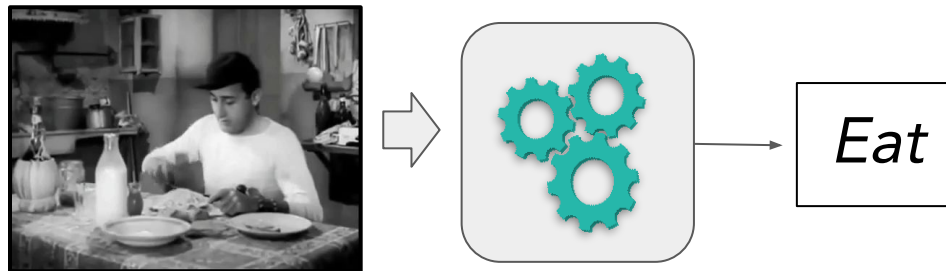
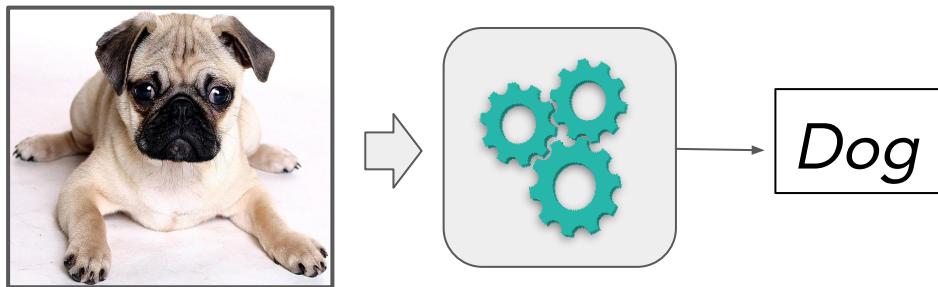
Paolo Rabino

paolo.rabino@polito.it

Overview

1. Train a **Convolutional LSTM** for First Person Action Recognition (FPAR):
 - Network: **ResNet 34 + RNN**
 - Videos: **GTEA61**
2. Exercise **Steps**:
 - Learning without Temporal Information (Avgpool)
 - Learning with Temporal Information (LSTM)
 - Learning with Spatio-Temporal Information (ConvLSTM)

Introduction



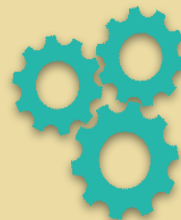
Introduction



Introduction




Network



Introduction



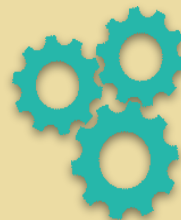
Network



Introduction




Network



Introduction




Network



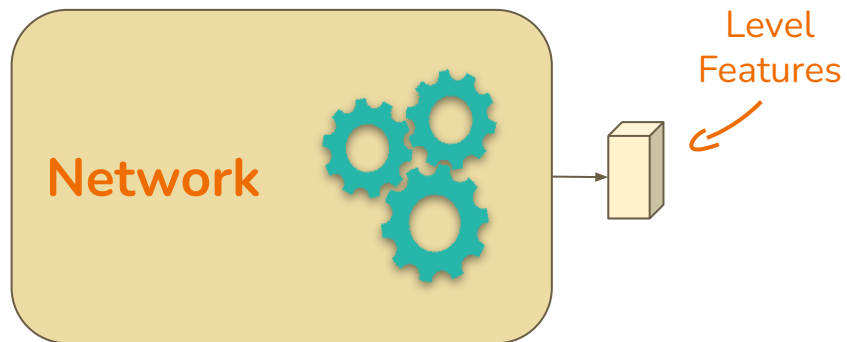
Introduction



Network



Introduction



Introduction

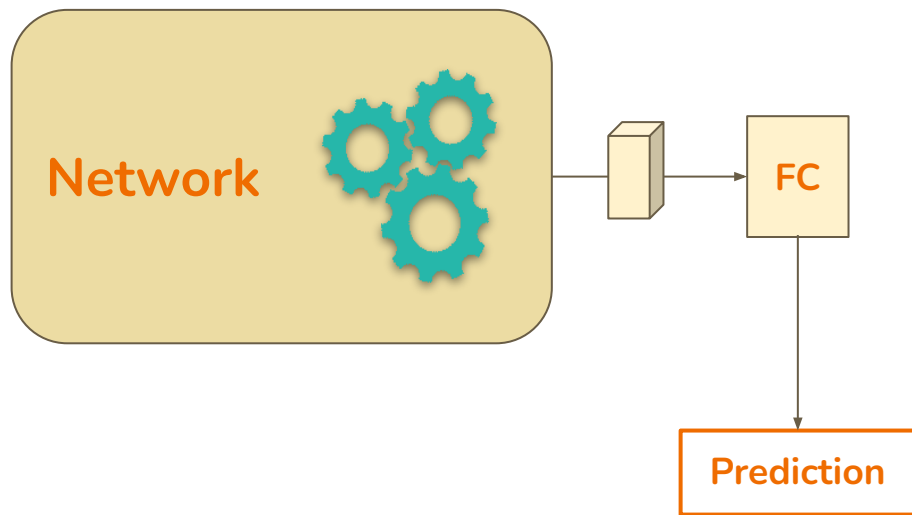


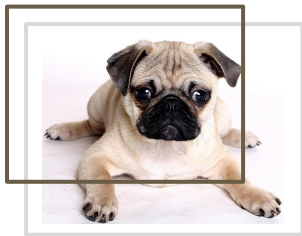
Image vs Frames

Image vs Frames

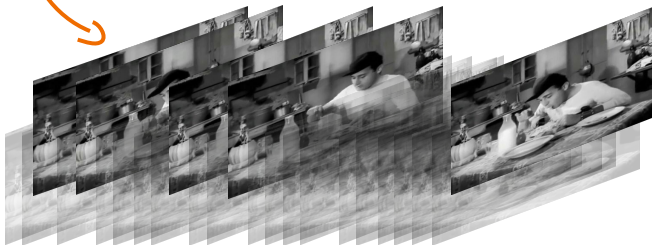
Don't worry!
Already
Implemented!!!

Training

Random Crop




Random Crop + **Random Sampling**



Test

Center Crop



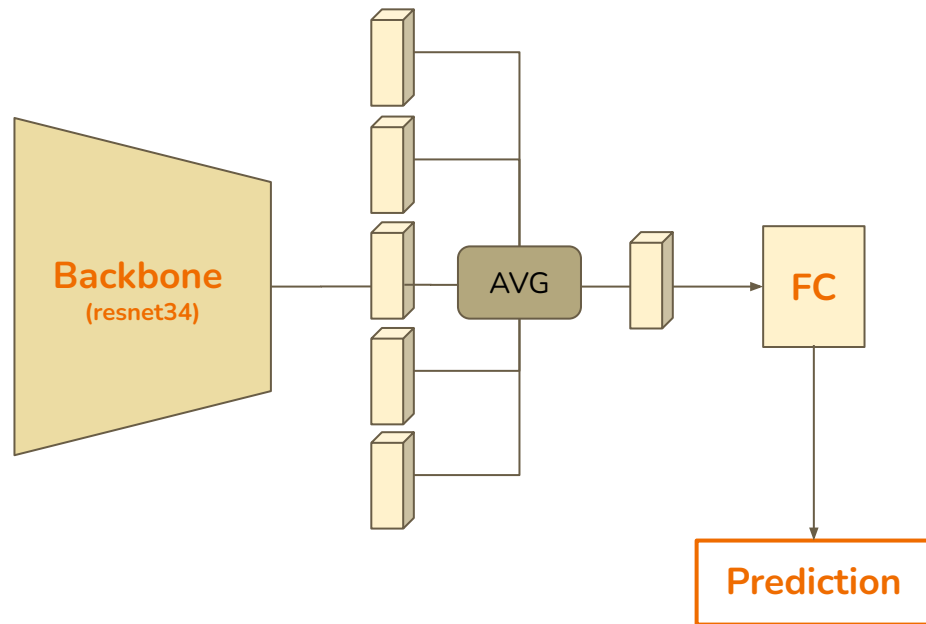
Center Crop + **Uniform Sampling**



Network - 0

Learning without Temporal Information (avg)


Network - 0



Network - 1

Learning with Temporal Information (LSTM)

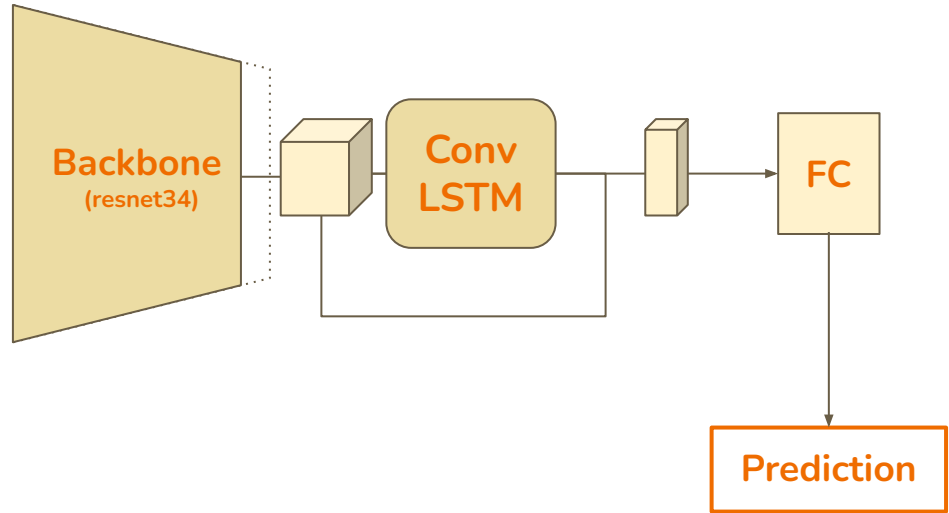
Network - 1



Network - 2

Learning with Spatio-Temporal Information (ConvLSTM)

Network - 2



Dataset: GTEA61

Dataset:GTEA61

- 61 actions
- 4 users:
 - S1, S2, S3, S4

Training sets = S1 S3 S4 (labeled)

Validation set = Test set = S2.



Dataset:

https://drive.google.com/drive/folders/1_NAcoR0UGH1eLsiWMOx_Py8yeAocknA2?usp=sharing

Code templates


The template of the main code is available [here](#):

Save a copy on your own drive:

“File” -> “Save a copy in Drive”

Switch to GPU acceleration:

“Edit” -> “Notebook Settings”



Step 0: Before you start

Learning without Temporal Information (avg)

Before you start

1. Study code and data:
 - a. Read carefully the template code (including the comments) to understand how everything is done.
 - b. Explore the data provided.

2. Run the code:
 - a. try to run the code, “Learning without Temporal Information” is already implemented
 - b. you have to stay connected

Step 1: LSTM

Learning with Temporal Information (LSTM)

Step 1: LSTM

1. Set the variable `homework_step = 1` in MAIN PARAMs
2. Implement the **LSTM** in the class `MyLSTM`, section `Model`
 - a. you should implement the model in the `init` function and the `forward`
3. Train only the **LSTM** and the **Classifier** (maintaining freezed the backbone)
 - a. follow the same procedure used in `Build Model - Loss - Opt` and in `Training` to train only the classifier in the Step 0

Hint: Use the LSTM with `frame_feat`

Step 2: ConvLSTM

Learning with Spatio-Temporal Information (ConvLSTM)

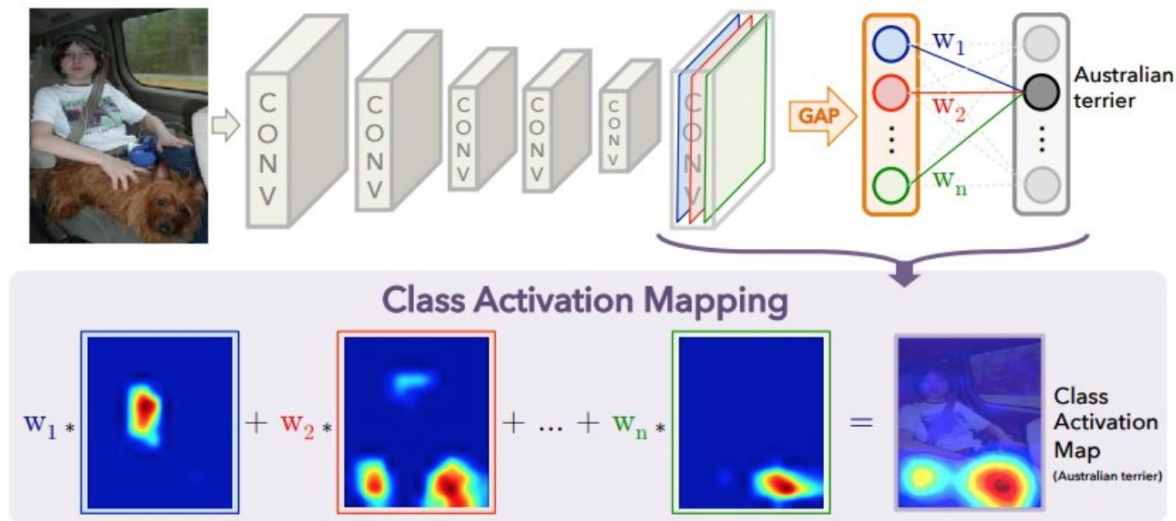
Step 2: ConvLSTM

1. Set the variable `homework_step = 2` in MAIN PARAMs
2. Implement the **ConvLSTM** in the class `MyConvLSTM`, section `Model`
 - a. it is very similar to the standard LSTM with the difference that it use the convolution operation instead of `nn.linear` and it works on `spatial_frames_feat` instead of `frames_feat`
3. Feed into the **ConvLSTM** the features before the avgpool of the resnet34
 - a. named `spatial_frame_feat`
4. Load the pretrained-weights that you find in
“`/content/best_model_state_dict_rgb_split2.pth`”
5. Train only the **ConvLSTM** and the **Classifier**, maintaining freezed the backbone

Step 4: Class Activation Map (CAM)

Extra!

Step 4: Class Activation Map (CAM)



Step 4: Class Activation Map (CAM)

```
logit, feature_conv, _ = self.resNet(inputVariable[t])
```

```
bz, nc, h, w = feature_conv.size()
```

```
feature_conv1 = feature_conv.view(bz, nc, h*w)
```

```
probs, idxs = logit.sort(1, True)
```

```
class_idx = idxs[:, 0]
```

```
cam =  
torch.bmm(self.weight_softmax[class_idx].unsqueeze(1),  
feature_conv1)
```

```
cam_img = F.softmax(cam, 1).data
```

```
cam_img = cam_img.cpu().numpy()
```

```
cam_img = cam_img.reshape(h, w)
```

```
cam_img = cam_img - np.min(cam_img)
```

```
cam_img = cam_img / np.max(cam_img)
```

```
cam_img = np.uint8(255 * cam_img)
```

```
output_cam = cv2.resize(cam_img, size_upsample)
```

```
img = cv2.cvtColor(np.uint8(img), cv2.COLOR_RGB2BGR)
```

```
heatmap = cv2.applyColorMap(output_cam, cv2.COLORMAP_JET)
```

```
result = heatmap * 0.3 + img * 0.5
```

NOW IT'S YOUR TURN, TRY!

Special thanks to Mirco Planamente for sharing the material

